

This is a copy of a shared ChatGPT chat. It won't be added to memory.

Report conversation

Regarding this question:

Then analyze one current OpenClaw or comparable agent-framework security advisory. In 500–700 words identify the vulnerable component, attacker-controlled input, preconditions, root cause, impact, fix, and one regression test. Explain whether prompt injection is required, merely helpful, or irrelevant.

I do not need you to write the security advisory. Instead could you please recommend a framework to analyze

Yes. For **this assignment**, I would recommend analyzing an **actual OpenClaw GitHub Security Advisory**, rather than a broad academic paper or a generic agent-framework vulnerability.

My recommendation: OpenClaw GHSA-2qj5-gwg2-xwc4

“OC-19: Unsanitized CWD path injection into LLM prompts” is a particularly good fit. It is a real OpenClaw advisory, **CVE-2026-27001**, affecting versions before 2026.2.15 .

 GitHub +1

Why I think it fits your assignment well:

Assignment requirement	Why this advisory fits
Vulnerable component	OpenClaw's workspace/CWD handling and prompt construction
Attacker-controlled input	A directory/workspace path containing control/format characters
Preconditions	Attacker needs to influence the directory in which OpenClaw operates
Root cause	Unsanitized filesystem data is inserted into an LLM system prompt
Impact	Prompt injection can alter agent behavior, potentially causing unintended tool use or disclosure

Assignment requirement	Why this advisory fits
Fix	Sanitize control/format characters before inserting the path into the prompt
Regression test	Create a maliciously named workspace and verify the resulting prompt cannot interpret the path as instructions
Prompt injection question	Essential to the exploit mechanism, rather than merely incidental

The advisory is also nice because the security boundary is easy to explain: **data that should be treated as a filesystem path crosses into the agent's instruction/prompt context without sufficient sanitization.**  [GitHub](#)

One important thing

There is an even more interesting OpenClaw advisory if your professor wants something focused on **agent security boundaries rather than prompt injection**:

CVE-2026-42434 — “Sandboxed agents could escape exec routing via host=node override.”

That vulnerability involved a sandboxed agent requesting `host: "node"` and causing execution to be routed to a remote node instead of remaining within the intended sandbox. It was fixed in OpenClaw `2026.4.10`.  [GitHub](#)

For your particular prompt, though, I'd choose **CVE-2026-27001** because it gives you a very clean answer to the professor's explicit question:

“Explain whether prompt injection is required, merely helpful, or irrelevant.”

You can directly analyze that distinction instead of forcing it into an unrelated vulnerability.

Suggested structure for your 500–700 words

I would organize your research notes around exactly these headings:

1. Vulnerable component
2. Attacker-controlled input
3. Preconditions
4. Root cause
5. Impact
6. Fix

7. Regression test

8. Role of prompt injection

That maps almost one-to-one with the assignment requirements and should make the final 500–700 words pretty easy to assemble.

The official advisory is here: [OpenClaw GHSA-2qj5-gwg2-xwc4](#) 

Also, OpenClaw's current security policy explicitly distinguishes **prompt-injection-only reports** from vulnerabilities that actually cross a policy, authorization, sandbox, or tool boundary, which is useful context when discussing why this particular issue qualifies as a security advisory.  [GitHub](#)

 Sources

ChatGPT can make mistakes. Check important info